

Midterm Exam

- Thursday, October 24, 7:15 – 9:15 p.m.
Last name A - Lin in room B130 Van Vleck
Last name Liou - Z in room 3650 Humanities
- Covers topics through Decision Trees, Random Forests, and k -Nearest-Neighbors
- Closed book except 8.5" x 11" sheet with notes on both sides (typed or handwritten)
- Bring student ID number, pencil, eraser, calculator (not on a phone)

1

- Covers lecture notes, readings in textbook (except Chapters 1 and 2), and 1 paper (intro to machine learning)
- True/False and multiple choice questions

2

Uninformed Search Methods

- Problem solving as search, problem representation in terms of states, goal test, operators, state-space graph search formulation, closed world assumption, Frontier and Explored lists, expanding a node, partial solution path, solution path, search tree
- Properties and computations
 - completeness, optimality, time and space complexity
- Methods
 - breadth-first search
 - depth-first search
 - uniform-cost search
 - iterative-deepening search
 - bidirectional search

3

Informed Search Methods

- Heuristic functions, evaluation function, admissible heuristic ($h \leq h^*$), consistent heuristic ($h(n) \leq c(n, n') + h(n')$), better informed heuristic, devising heuristics, completeness, admissibility, optimality
- Methods
 - best-first search
 - greedy best-first search
 - beam search
 - algorithm A
 - algorithm A*
 - IDA*

4

Local Search Methods

- Local search problem formulation, operators, neighborhood (aka move set), local optima problem
- Methods:
 - hill-climbing algorithm
 - hill-climbing with random restarts,
 - Simulated annealing (stochastic hill-climbing) escaping local optima, Boltzman's equation ($p = e^{\Delta E / T}$), temperature, cooling schedule

5

Game Playing

- Zero-sum games, perfect information games, deterministic vs. stochastic games, game playing as search, search tree, branching factor, ply, static evaluation function, horizon effect
- Methods
 - Minimax algorithm
 - Minimax principle, optimal playing strategy
 - Alpha-beta pruning algorithm
 - Cutoff tests: If $v \leq \alpha$ for *some* MAX node ancestor, **don't visit** any more of the current MIN node's children. If $v \geq \beta$ for *some* MIN node ancestor, **don't visit** any more of the current MAX node's children
 - iterative-deepening search with alpha-beta pruning
 - non-deterministic games
 - chance nodes, expectiminimax value
 - Monte Carlo tree search
 - Pure MCTS, Selection, expansion, simulation, backpropagation, payout, exploitation, exploration

7

MCTS Algorithm

Recursively build search tree, where each iteration consists of:

1. Selection: Starting at root, successively select best child nodes using scoring method until leaf node L reached
2. Expansion: Create and add best (or random) new child node, C , of L
3. Simulation: Perform a (random) payout from C
4. Backpropagation: Update score at C and all of C 's ancestors in search tree based on payout results

8

Unsupervised Learning

- Inductive learning problem, feature space, feature, attribute, examples (aka instances), labels, classes, training set, tuning set, testing set, classification problems, decision boundaries
- Methods
 - Hierarchical Agglomerative Clustering (HAC)
 - single linkage, complete linkage, average linkage, Euclidean distance, Manhattan distance, dendrogram
 - k -Means Clustering
 - cluster center (centroid), distortion cluster quality measure: For all clusters, sum of squared distances from each point to its cluster center
 - Nothing on mean-shift clustering

9

Hierarchical Agglomerative Clustering Algorithm

Input: a training sample $\{x_i\}_{i=1}^n$; a distance function $d()$.

1. Initially, place each instance in its own cluster (called a singleton cluster).
2. **while** (number of clusters > 1) **do**:
3. Find the closest cluster pair A, B , i.e., they minimize $d(A, B)$.
4. Merge A, B to form a new cluster.

Output: a binary tree showing how clusters are gradually merged from singletons to a root cluster, which contains the whole training sample.

10

K-Means Clustering Algorithm

- Input: x_1, \dots, x_n, k where each x_i is a point in a d -dimensional feature space
- **Step 1:** Select k cluster centers, c_1, \dots, c_k
- **Step 2:** For each point x_i , determine its cluster: Find the closest center (using, say, Euclidean distance)
- **Step 3:** Update all cluster centers as the centroids

$$c_i = \frac{1}{\text{num_pts_in_cluster_}i} \sum_{x \in \text{cluster } i} x$$

- Repeat steps 2 and 3 until cluster centers no longer change

11

Supervised Learning

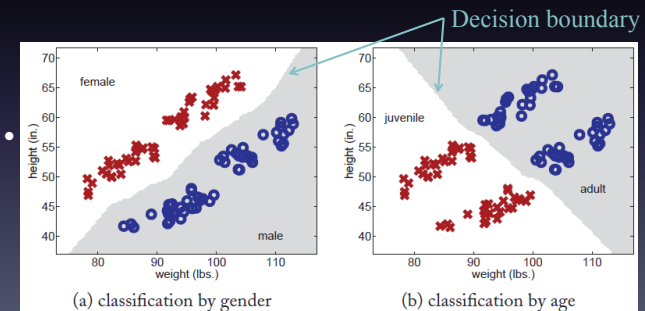
- Inductive bias, preference bias, decision boundaries, training set, tuning set, testing set, overfitting, noisy data, setting parameters
- Methods
 - K-nearest neighbor (k-NN)
 - Decision Trees
 - Ockham's razor, entropy, conditional entropy (aka remainder), information gain, categorical attributes, real-valued attributes, rule extraction, methods to avoid overfitting, pruning algorithm
 - Random Forests
 - Ensemble methods, bagging, bootstrap sampling, randomized node optimization, majority/mode classifier
 - Performance evaluation
 - Training set accuracy, training set error, testing set accuracy, testing set error, k -fold cross validation, leave-one-out cross validation

12

k-Nearest-Neighbors (k-NN)

Input: Training data $(x_1, y_1), \dots, (x_n, y_n)$; distance function $d()$; number of neighbors k ; test instance x^*

1. Find the k training instances x_{i_1}, \dots, x_{i_k} closest to x^* under distance $d()$.
2. Output y^* as the majority class of y_{i_1}, \dots, y_{i_k} . Break ties randomly.



13

- Entropy:

$$H(Y) = \sum_{i=1}^k -p_i \log_2 p_i$$

- Conditional Entropy:

$$H(Y | X = v) = \sum_{i=1}^k -\Pr(Y = y_i | X = v) \log_2 \Pr(Y = y_i | X = v)$$

$$H(Y | X) = \sum_{v: \text{values of } X} \Pr(X = v) H(Y | X = v)$$

- Information Gain:

$$I(Y; X) = H(Y) - H(Y | X)$$

14

Decision-Tree-Learning Algorithm

```

buildtree(examples, attributes, default-label)
  if empty(examples) then return default-label
  if (examples all have same label y) then return y
  if empty(attributes) then return majority-class of examples
  q = best_attribute(examples, attributes)
  tree = create-node with attribute q
  foreach value v of attribute q do
    v-ex = subset of examples with q == v
    subtree = buildtree(v-ex, attributes - {q}, majority-class(examples))
    add arc from tree to subtree
  return tree

```

15

Pruning using a Greedy Algorithm

Prune(tree T, TUNE set)

1. Compute T's accuracy on TUNE; call it Acc(T)
2. For every internal node N in T:
 - a) New tree T_N = copy of T, but prune (delete) the subtree under N
 - b) N becomes a leaf node in T_N . The class is the majority vote of TRAIN examples reaching N
 - c) $\text{Acc}(T_N)$ = T_N 's accuracy on TUNE
3. Let T^* be the tree (among the T_N 's and T) with the largest Acc()
Set $T = T^*$ /* prune */
4. If no improvement then Return T else Goto Step 1

16

Random Forests

For each tree,

1. Build a training set by choosing n times with replacement from all N available training examples (aka "taking a bootstrap sample")
2. At each node of decision tree during construction, choose a *random subset* of m **attributes** from the total number, M , of possible attributes ($m \ll M$)
3. Select the best attribute at node using Max-Gain

17

K-Fold Cross Validation

1. Divide all examples into K disjoint subsets
 $E = E_1, E_2, \dots, E_K$
2. For each $i = 1, \dots, K$
 - let TEST set = E_i and TRAIN set = $E - E_i$
 - build decision tree using TRAIN set
 - determine accuracy Acc_i using TEST set
3. Compute **K-fold cross-validation** estimate of performance = mean accuracy =
 $(Acc_1 + Acc_2 + \dots + Acc_K)/K$

18

Leave-One-Out Cross Validation

For $i = 1$ to N do // N = number of examples

1. Let (x_i, y_i) be the i^{th} example
2. Remove (x_i, y_i) from the dataset
3. Train on the remaining $N-1$ examples
4. Compute accuracy on i^{th} example

- Accuracy = mean accuracy on all N runs

19