

Constraint Satisfaction Problems


Chapter 6.1 – 6.4

Derived from slides by S. Russell and P. Norvig, A. Moore, and R. Khoury

Constraint Satisfaction Problems (CSPs)


- Standard search problem:
 - **state** is a "black box" – any data structure that supports successor function, heuristic function, and goal test
- CSP:
 - **state** is defined by **variables** X_i with **values** from **domain** D_i
 - **goal test** is a set of **constraints** specifying allowable combinations of values for subsets of variables
 - Use *variable-based model*
 - Solution is not a path but an *assignment of values for a set of variables that satisfy all constraints*

Example: 8-Queens




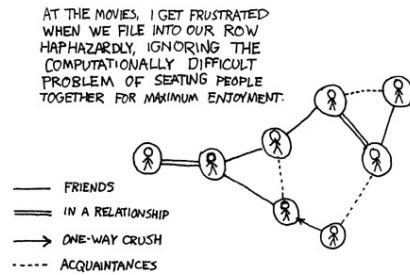
Example: Cryptarithmic

$$\begin{array}{r} \text{ T W O} \\ + \text{ T W O} \\ \hline \text{ F O U R} \end{array}$$




- **Variables:** $F, T, U, W, R, O, X_1, X_2, X_3$
- **Domains:** $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- **Constraints:** $\text{Alldiff}(F, T, U, W, R, O)$
 - $O + O = R + 10 \cdot X_1$
 - $X_1 + W + W = U + 10 \cdot X_2$
 - $X_2 + T + T = O + 10 \cdot X_3$
 - $X_3 = F, T \neq 0, F \neq 0$

Example: Movie Seating



Example: Graph Coloring




- Each circle marked $V_1 \dots V_6$ we must assign R, G or B
- No two adjacent circles may be assigned the same value
- Note: 2 variables have already been assigned a color in this example

Other Applications of CSPs

- Assignment problems
 - e.g., who teaches what class
- Timetable problems
 - e.g., which class is offered when and where?
- Scheduling problems
- VLSI or PCB layout problems
- Boolean satisfiability
- N-Queens
- Graph coloring
- Games: Minesweeper, Magic Squares, Sudoku, Crosswords
- Line-drawing labeling

Note: many problems require real-valued variables


Example: Map-Coloring



- **Variables:** WA, NT, Q, NSW, V, SA, T
- **Domains:** $D_i = \{\text{red, green, blue}\}$
- **Constraints:** adjacent regions must have different colors
 e.g., $WA \neq NT$, or $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), (\text{blue, red}), (\text{blue, green})\}$

Note: In general, 4 colors are necessary


Example: Map-Coloring



Solutions are **complete** (i.e., all variables are assigned values) and **consistent** (i.e., does not violate any constraints) assignments, e.g., WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green

Constraint Graph

- **Binary CSP**: each constraint relates **two** variables
- **Constraint graph**: nodes are **variables**, arcs are **constraints**



Varieties of CSPs

- Discrete variables
 - finite domains:
 - n variables, domain size $d \rightarrow O(d^n)$ complete assignments
 - e.g., Boolean CSPs, Boolean satisfiability
 - infinite domains:
 - integers, strings, etc.
 - e.g., job scheduling, variables are start/end times for each job
- Continuous variables
 - e.g., start/end times for Hubble Space Telescope observations
 - linear constraints solvable in polynomial time by linear programming

Kinds of Constraints

- **Unary** constraints involve a single variable
 - e.g., SA \neq green
- **Binary** constraints involve pairs of variables
 - e.g., SA \neq WA
- **Higher-order** constraints involve 3 or more variables
 - e.g., cryptarithmic column constraints

Local Search for CSPs

- Hill-climbing, simulated annealing, genetic algorithms typically work with "complete" states, i.e., *all* variables have values at every step
- To apply to CSPs:
 - allow states with some *unsatisfied* constraints
 - operators **assign** a value to a variable
- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts heuristic**:
 - choose value that *violates the fewest constraints*, i.e., hill-climb by minimizing $f(n)$ = total number of violated constraints


Local Search

Min-Conflicts Algorithm:


- Assign to each variable a random value, defining the initial state
- while** state not consistent **do**
 - Pick a variable, *var*, that has constraint(s) violated
 - Find value, *v*, for *var* that minimizes the *total* number of violated constraints (over all variables)
 - var* = *v*

Example: 4-Queens

- States**: 4 queens in 4 columns ($4^4 = 256$ states)
- Actions**: move queen to new row in its column
- Goal test**: no attacks
- Evaluation function**: $f(n)$ = total number of attacks



Min-Conflicts Algorithm



Min-Conflicts Algorithm


- Advantages
 - Simple and Fast: Given random initial state, can solve n -Queens in almost constant time for arbitrary n with high probability (e.g., $n = 1,000,000$ can be solved on average in about 50 steps!)
- Disadvantages
 - Only searches states that are reachable from the initial state
 - Might not search entire state space
 - Does not allow worse moves (but can move to a neighbor with the *same* cost)
 - Might get stuck in a local optimum
 - Not complete
 - Might not find a solution even if one exists

Standard Tree Search Formulation

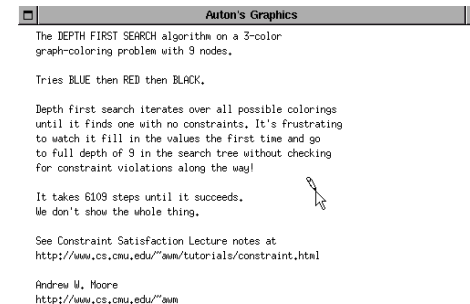
States are defined by all the values *assigned so far*

- **Initial state:** the empty assignment $\{ \}$
- **Successor function:** assign a value to an unassigned variable
- **Goal test:** the current assignment is *complete* and *consistent*, i.e., all variables assigned a value and all constraints satisfied
- Goal: Find **any** solution, so cost is not important
- Every solution appears at depth n with n variables
→ use depth-first search

DFS for CSPs



- Variable assignments are **commutative**, i.e.,
[WA=R then NT=G] same as [NT=G then WA=R]
- What happens if we do DFS with the order of assignments as *B* tried first, then *G*, then *R*?
- **Generate-and-test strategy:** Generate candidate solution, then test if it satisfies all the constraints
- This makes DFS look very stupid!
- Example:
<http://www.cs.cmu.edu/~awm/animations/constraint/9d.html>



Improved DFS: Backtracking w/ Consistency Checking

- Don't generate a successor that creates an inconsistency with any *existing* assignment, i.e., perform **consistency checking** *when node is generated*
- Successor function assigns a value to an unassigned variable that does **not** conflict with *all* current assignments
 - Deadend if no legal assignments (i.e., no successors)
- Backtracking (DFS) search is the basic uninformed algorithm for CSPs
- Can solve n -Queens for $n \approx 25$

Backtracking w/ Consistency Checking

Start with empty state

while not all vars in state assigned a value **do**

 Pick a variable (randomly or with a heuristic)

if it has a value that does not violate any constraints

then Assign that value


else

 Go back to previous variable and assign it another value


Backtracking Example




Australia Constraint Graph




Backtracking Example



Backtracking Example

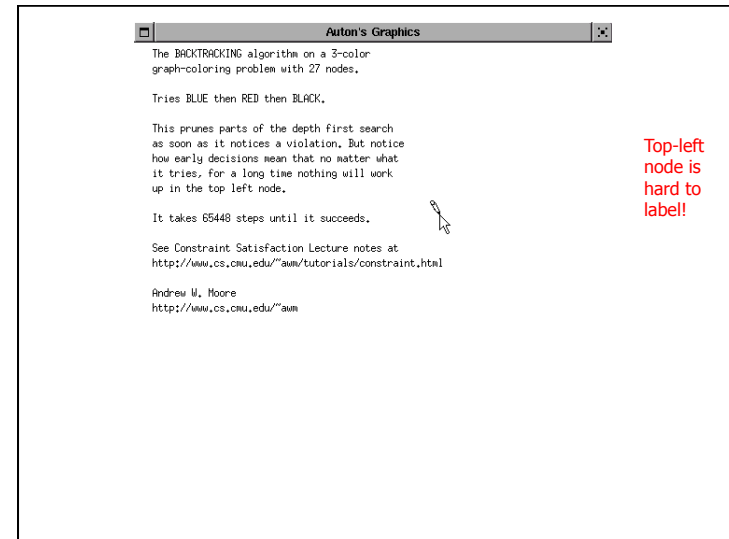
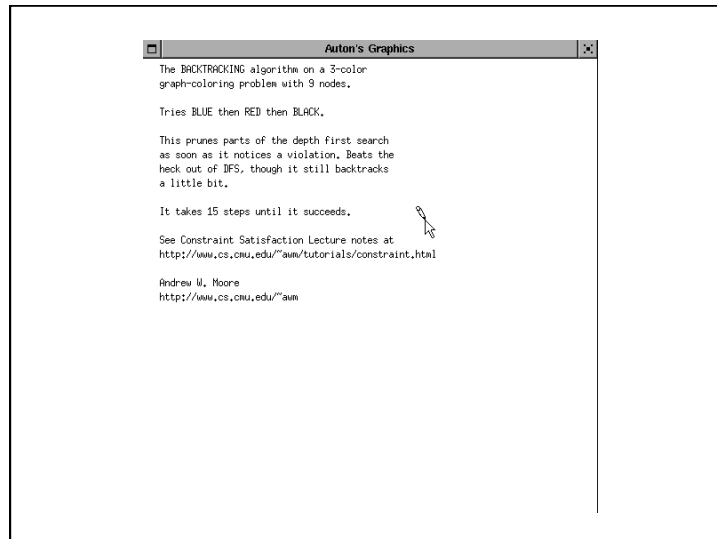


Backtracking Example



Backtracking Search

- Depth-first search algorithm
 - Goes down one variable at a time
 - At a deadend, backs up to *last* variable whose value can be changed without violating any constraints, and changes it
 - If you backed up to the root and tried all values, then there is *no* solution
- Algorithm is *complete*
 - Will find a solution if one exists
 - Will expand the entire (finite) search space if necessary
- Depth-limited search with depth limit = n



Improving Backtracking Efficiency

- **Heuristics** can give huge gains in speed
 - Which *variable* should be assigned next?
 - In what order should its *values* be tried?
 - Can we detect inevitable failure early?

Which Variable Next? Most-Constrained Variable


- **Most-constrained variable**
 - Choose the variable with the *fewest* legal values



- Called the **minimum remaining values (MRV)** heuristic
- Minimizes branching factor; Tries to cut off search ASAP


Which Variable Next? Most-Constraining Variable

- Tie-breaker among most-constrained variables
- **Most-constraining variable**
 - Choose the variable with the *most* constraints on the *remaining* variables
- Called the **degree heuristic**
- Tries to cut off search ASAP




Which Value Next? Least-Constraining Value

- Given a variable, choose the **least-constraining value**:
 - i.e., the one that rules out the *fewest* values in the remaining variables
 - try to pick values *best first*



- Combining these heuristics makes 1000-Queens feasible


Forward Checking Algorithm



- Initially, for each variable, record the **set of all possible legal values for it**
- When you assign a value to a variable in the search, *update the set of legal values for all unassigned variables. Backtrack immediately if you empty a variable's set of possible values.*

Forward Checking Algorithm

- Keep track of **remaining legal values** for all variables
- Deadend when any variable has **no** legal values




Example: Map-Coloring



- **Variables:** WA, NT, Q, NSW, V, SA, T
- **Domains:** $D_i = \{\text{red, green, blue}\}$
- **Constraints:** adjacent regions must have different colors
e.g., $WA \neq NT$, or $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), (\text{blue, red}), (\text{blue, green})\}$


Constraint Graph

- **Binary CSP:** each constraint relates **two** variables
- **Constraint graph:** nodes are **variables**, arcs are **constraints**



Forward Checking

- Keep track of **remaining legal values** for all unassigned variables
- Deadend when any variable has **no** legal values



Note: WA is **not** the most constraining var

WA	NT	Q	NSW	V	SA	T
red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue
red	green	blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue

Forward Checking

- Keep track of **remaining legal values** for all unassigned variables
- Deadend when any variable has **no** legal values




Note: Q is **not** most constrained variable

WA	NT	Q	NSW	V	SA	T
red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue
red	green	blue	red, green, blue	red, green, blue	red, green, blue	red, green, blue
red	green	blue	red	red, green, blue	red, green, blue	red, green, blue

Forward Checking

- Keep track of **remaining legal values** for all unassigned variables
- Deadend when any variable has **no** legal values



Auton's Graphics

The FORWARD CHECKING algorithm on a 3-color graph-coloring problem with 27 nodes.

Tries BLUE then RED then BLACK.

Little dots denote the availability lists for the nodes.


Notice that unlike backtracking search, Forward Checking realizes as soon as it tries setting the node at (row=bottom+1,col=rightmost-1) to Black that it's not going to be able to satisfy the top-left node.

See Constraint Satisfaction Lecture notes at <http://www.cs.cmu.edu/~aam/tutorials/constraint.html>

Andrew W. Moore
<http://www.cs.cmu.edu/~aam>


Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for **all** failures:



- NT and SA cannot *both* be blue!
- **Constraint propagation** repeatedly (recursively) enforces constraints for *all* variables

Constraint Propagation




Main idea: When you delete a value from a variable's domain, check all variables connected to *it*. If any of them change, delete all inconsistent values connected to *them*, etc.

Note: In the above example, nothing changes


Arc Consistency

- Simplest form of propagation makes each **arc** (i.e., each binary constraint) **consistent**
- $X \rightarrow Y$ is consistent if
for **every** value x at var X there is **some** allowed y ,
i.e., there is at least 1 value of Y that is consistent with x at X




Arc Consistency

- Simplest form of propagation makes each **arc** **consistent**
- $X \rightarrow Y$ is consistent if
for **every** value x at X there is **some** allowed y ;
if not, delete x



Arc Consistency


- Simplest form of propagation makes each **arc** **consistent**
- $X \rightarrow Y$ is consistent if
for **every** value x at X there is **some** allowed y ; if not,
delete x



- If X loses a value, **all** neighbors of X must be rechecked

Arc Consistency

- Simplest form of propagation makes each **arc** **consistent**
- $X \rightarrow Y$ is consistent if
for **every** value x at X there is **some** allowed y ; if not, delete x



- If X loses a value, all neighbors of X must be rechecked
- Arc consistency detects failure *earlier* than forward checking
- Use as a preprocessor and after each assignment during search

Auton's Graphics

The CONSTRAINT PROPAGATION algorithm on a 3-color graph-coloring problem with 27 nodes.

Tries BLUE then RED then BLACK.

Little dots denote the availability lists for the nodes.

Notice that unlike forward checking search, Constraint Propagation realizes very early on (on its third step) that (row=bottom+1,col=rightmost-1) must not be black and so (row=bottom,col=4) must not be red. It does better than forward checking and MUCH better than backtracking!

See Constraint Satisfaction Lecture notes at <http://www.cs.cmu.edu/~awm/tutorials/constraint.html>

Andrew W. Moore
<http://www.cs.cmu.edu/~awm>

Bottom-right node must *not* be red because node to upper-left must *not* be black

Arc Consistency Algorithm “AC-3”

```

function AC-3(csp) // returns false if inconsistency is found and
                        true otherwise
// input: csp, a binary CSP with components ( $X, D, C$ )
// local variables: queue, a queue of arcs; initially all arcs in csp
while queue not empty do {
    ( $X_i, X_j$ ) = Remove-First(queue); // check if  $X_i \rightarrow X_j$  consistent
    if Revise(csp,  $X_i, X_j$ ) then { // make arc consistent
        if size of  $D_j = 0$  then return false
        foreach  $X_k$  in  $X_i$ .Neighbors - { $X_j$ } do // propagate changes
                                                to neighbors
            add ( $X_k, X_i$ ) to queue
    }
}
return true

```


Arc Consistency Algorithm “AC-3”

```

function Revise(csp,  $X_i, X_j$ ) // returns true if we revise the
                                domain of  $X_j$ 
    revised = false;
    foreach  $x$  in  $D_j$  do { // check if  $X_i \rightarrow X_j$  consistent
        if no value  $y$  in  $D_j$  allows ( $x, y$ ) to satisfy the constraints
            between  $X_i$  and  $X_j$  then {
                delete  $x$  from  $D_j$ ;
                revised = true;
            }
    }
    return revised

```

Constraint Propagation



- In this example, constraint propagation solves the problem without search ... **But not always that lucky!**
- Constraint propagation can be done as a **preprocessing step**
- And it can be performed **during** search
 - Note: when you backtrack, you must *undo* some of your additional constraints

Combining Search with CSP

- Idea: Interleave search and CSP inference
- Perform DFS
 - At each node assign a selected value to a selected variable
 - Run CSP to reduce variables' domains and check if any inconsistencies arise as a result of this assignment

Combining Backtracking Search with CSP

```
function BACKTRACKING-SEARCH(csp) returns a solution or failure
return BACKTRACK({ }, csp)

function BACKTRACK(assignment, csp) returns a solution or failure
if assignment is complete then return assignment;
var = SELECT-UNASSIGNED-VARIABLE(csp);
foreach value in ORDER-DOMAIN-VALUES(var, assignment, csp) do {
  if value is consistent with assignment then {
    add {var = value} to assignment;
    inferences = AC-3(csp, var, value);
    if inferences != failure then {
      add inferences to assignment;
      result = BACKTRACK(assignment, csp);
      if result != failure then return result; }
  }
  remove {var = value} and inferences from assignment;
}
return failure
```

Summary

- CSPs are a special kind of problem:
 - states defined by values of a fixed set of variables
 - goal test defined by constraints on variable values
- Backtracking = depth-first search with one variable assigned per node plus consistency checking
- Variable ordering and value selection heuristics help significantly
- **Forward checking** prevents assignments that guarantee later failure
- Constraint propagation (e.g., **arc consistency**) does additional work to constrain values and detect inconsistencies earlier
- Iterative **min-conflicts** is usually effective in practice